



flickr

Computer vision at scale with Hadoop and Storm

Presented by: Huy Nguyen @huyng

Outline

- Computer vision at 10,000 feet
- Hadoop training system
- Storm continuous stream processing
- Hadoop batch processing
- Q&A

500+ million images are uploaded to
internet photo services every day

This upload rate is expected to
double by the end of 2014

- Mary Meeker (KPCB)
2013 Internet Trends

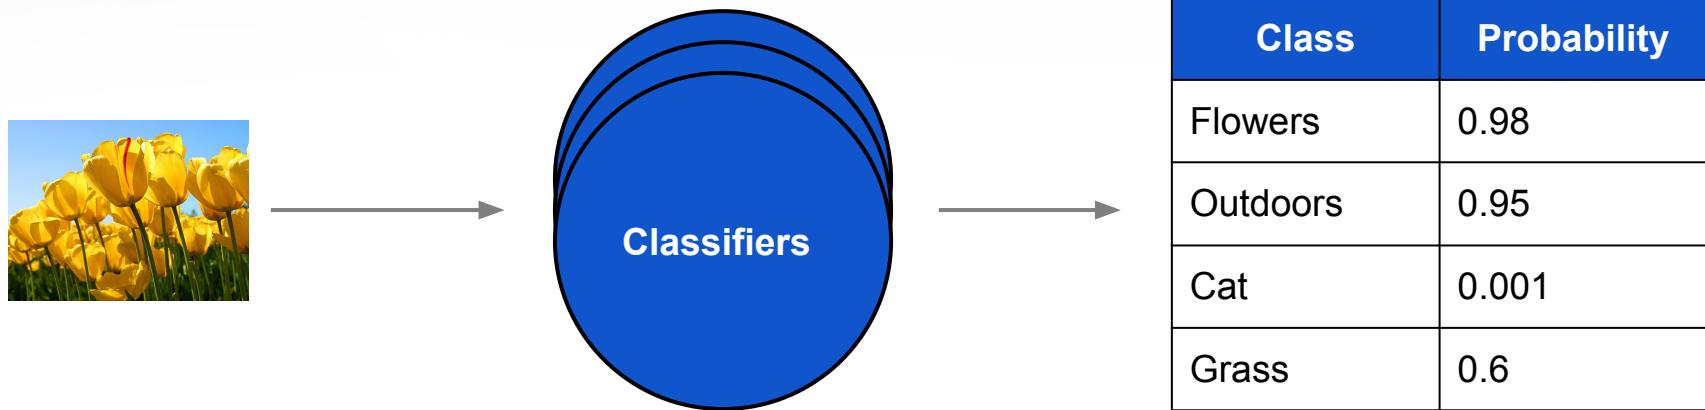


How can we help Flickr users better **organize, search, and browse** their photo collection?



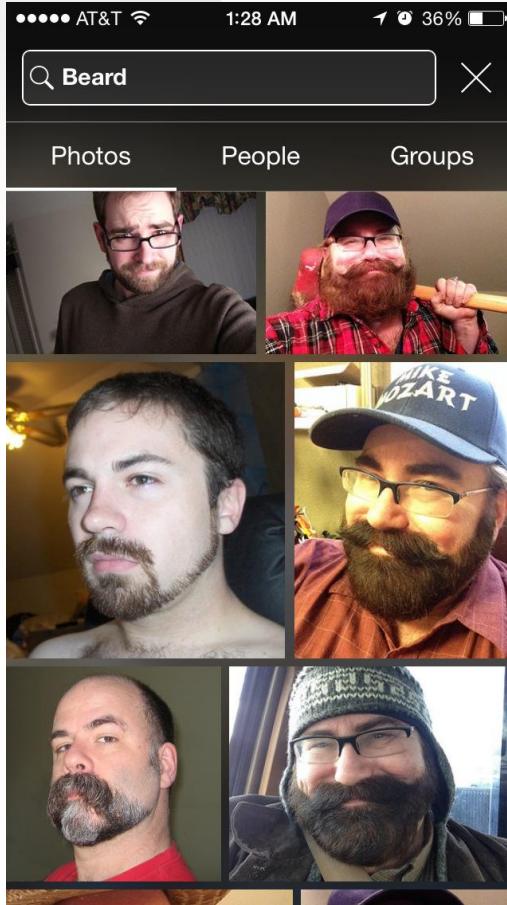
photo credit: Quinn Dombrowski

What is AutoTagging?



Any photo you upload to Flickr is now automatically classified using computer vision software

Demo



Auto Tags Feeding Search

- Flowers
- Buildings
- Outdoors
- Beach
- Beards
- and more!

Our Goals

- Train thousands of classifiers
- Classify millions of photos per day
- Compute auto tags on backlog of billions of photos
- Enable reprocessing bi-weekly and every quarter
- Do it all within 90 days of being acquired!

Our system for training classifiers

- High level overview of image classification
- Hadoop workflow

Classification in a nutshell

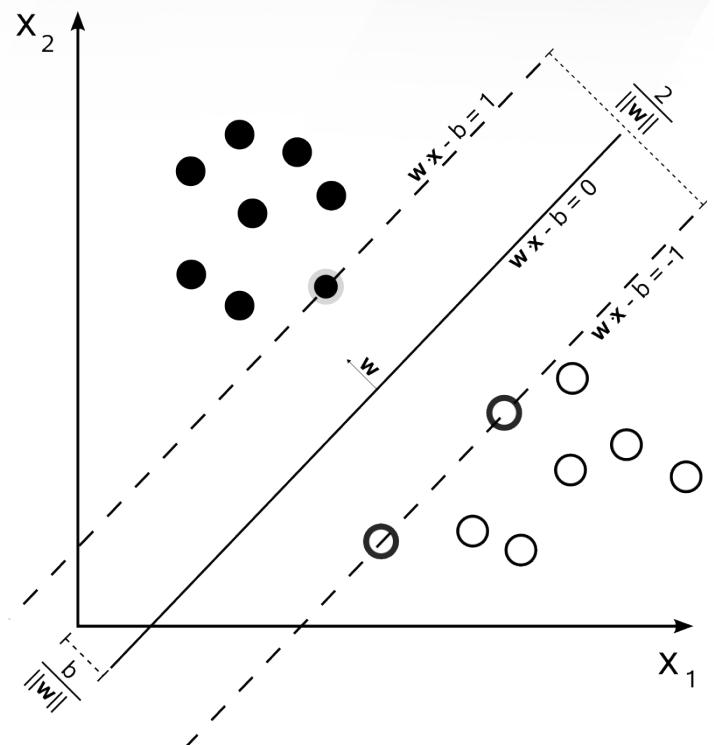
How do we define a classifier?

A classifier is simply the line “ $z = \mathbf{w}x - b$ ” which divides positive and negative examples with the largest margin.

Entirely defined by **w** and **b**

Training is simply adjusting **w** and **b** to find the line with the largest margin dividing negative and positive example data

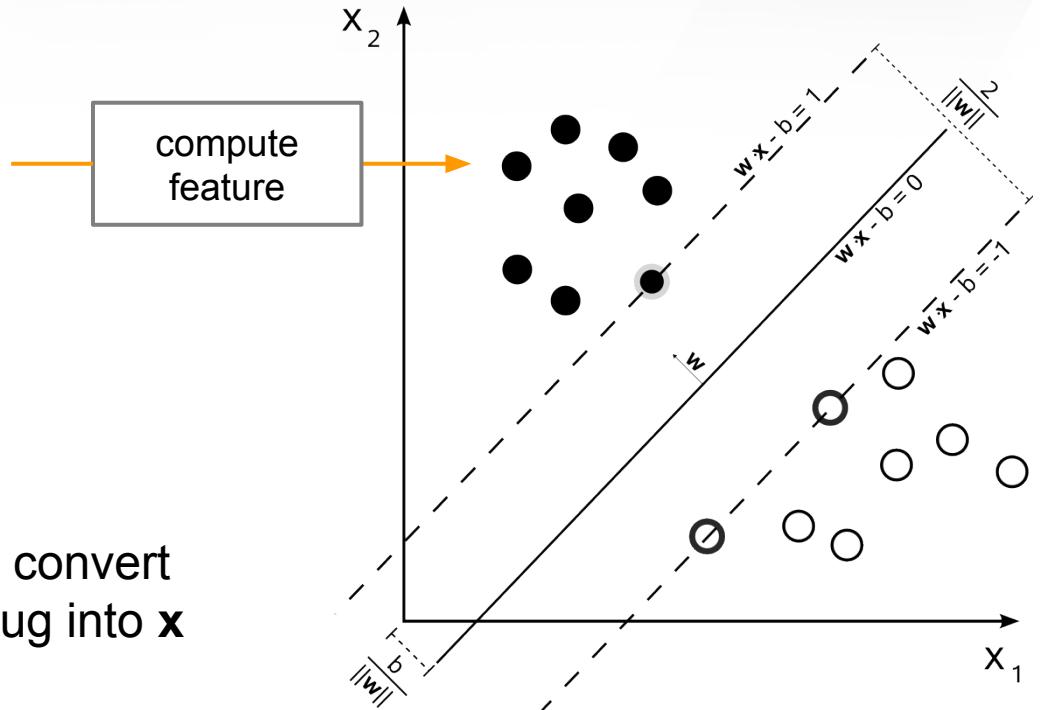
Classification is done by computing z, given some data point **x**



hadoop

YAHOO!

Image classification in a nutshell



To classify images we need to convert raw pixels into a “feature” to plug into \mathbf{x}



hadoop

YAHOO!

How are image classifiers made?



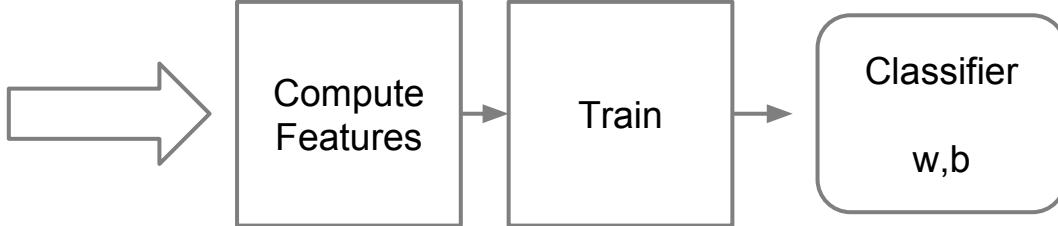
Class Name



Negative Examples
(approx. 100K)

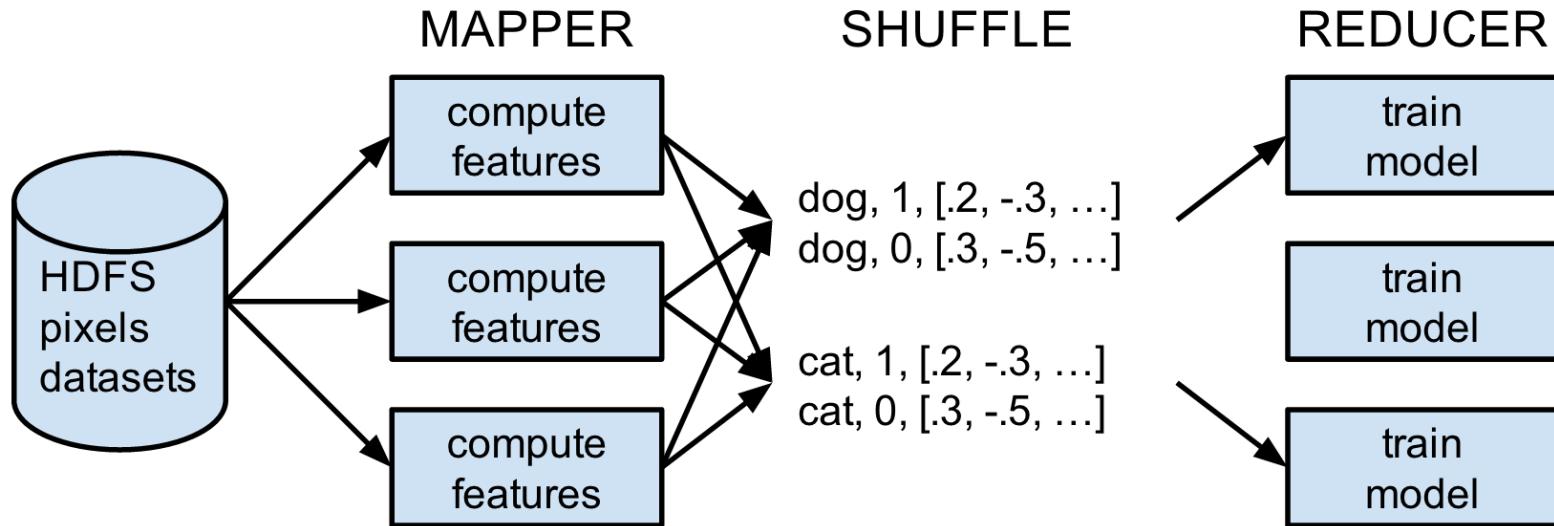


Positive Examples
(approx. 10K)

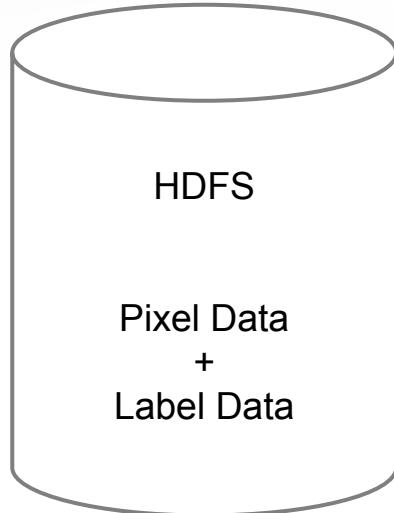


**Repeat a few thousand times
for each class**

Hadoop Classifier Training



Hadoop Classifier Training



Pixel Data

1	<base64>
2	<base64>
3	<base64>
4	<base64>
.	
.	
n	<base64>

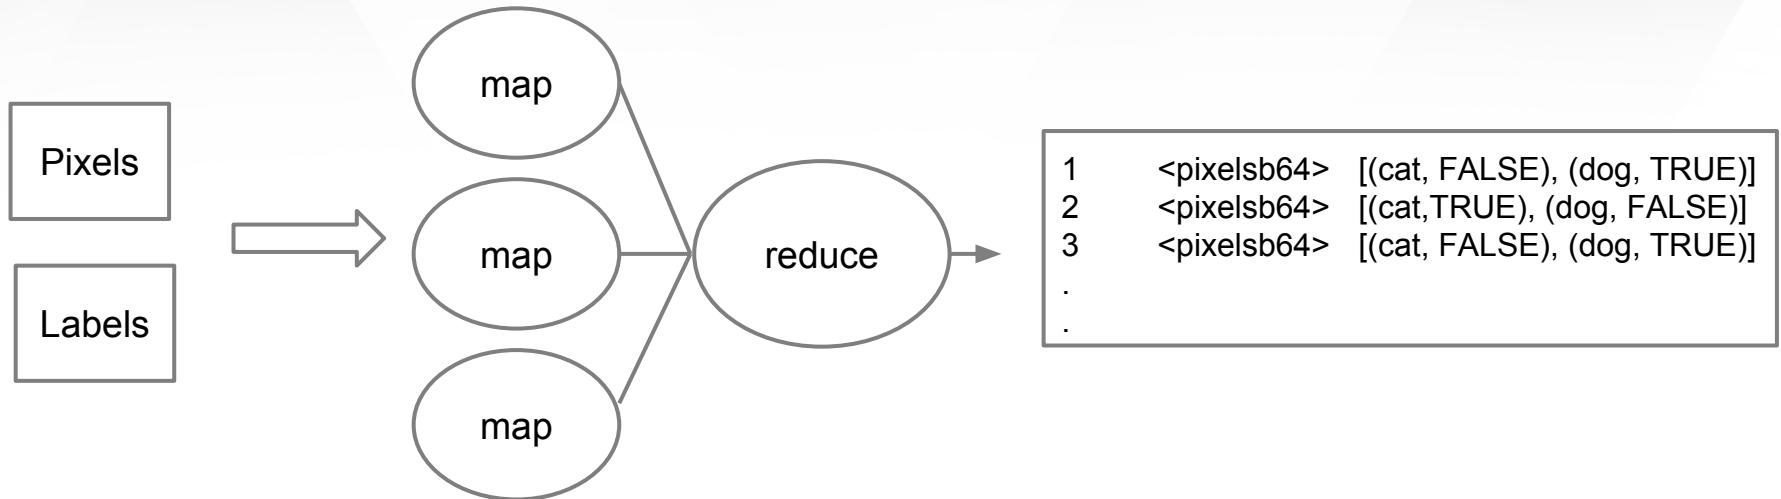
photo_id, pixel_data

Label Data

dog	1	TRUE
dog	2	FALSE
dog	3	TRUE
cat	1	FALSE
cat	2	TRUE
cat	3	FALSE
.		
.		

class, photo_id, label

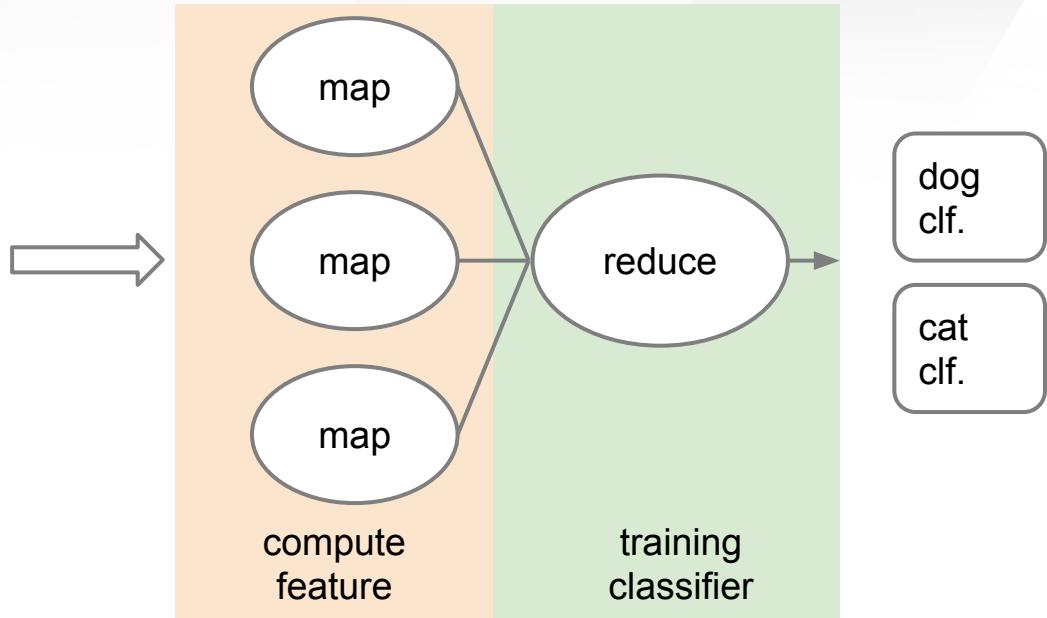
Hadoop Classifier Training - JOIN



- mappers emit (photo_id, pixels) **OR** (photo_id, class, label)
- reducers, takes advantage of sorting to emit (photo_id, pixels, json(label data))

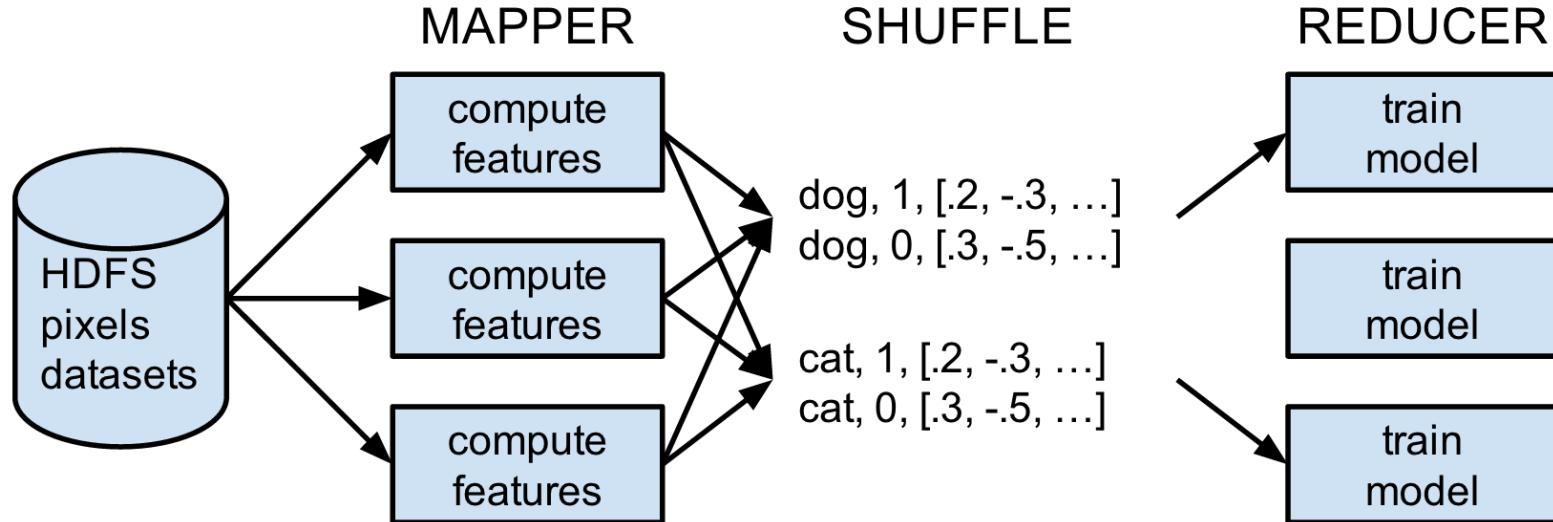
Hadoop Classifier Training - TRAIN

```
1 <pixelsb64> [(cat, FALSE), (dog, TRUE)]  
2 <pixelsb64> [(cat, TRUE), (dog, FALSE)]  
3 <pixelsb64> [(cat, FALSE), (dog, TRUE)]  
. .
```



- mappers compute features, emits a (class, feat64, label) **for each class image associated with image**
- reducers, takes advantage of sorting to train. Emits (class, w, b)

Hadoop Classifier Training



- Workflow coordinated with Oozie

Hadoop Classifier Training

- 10,000 mappers
- 1,000 reducers
- A few thousand classifiers in less than 6 hours
- Replaced OpenMPI based system
 - Easier deployment
 - Access to more hardware
 - Much simpler code

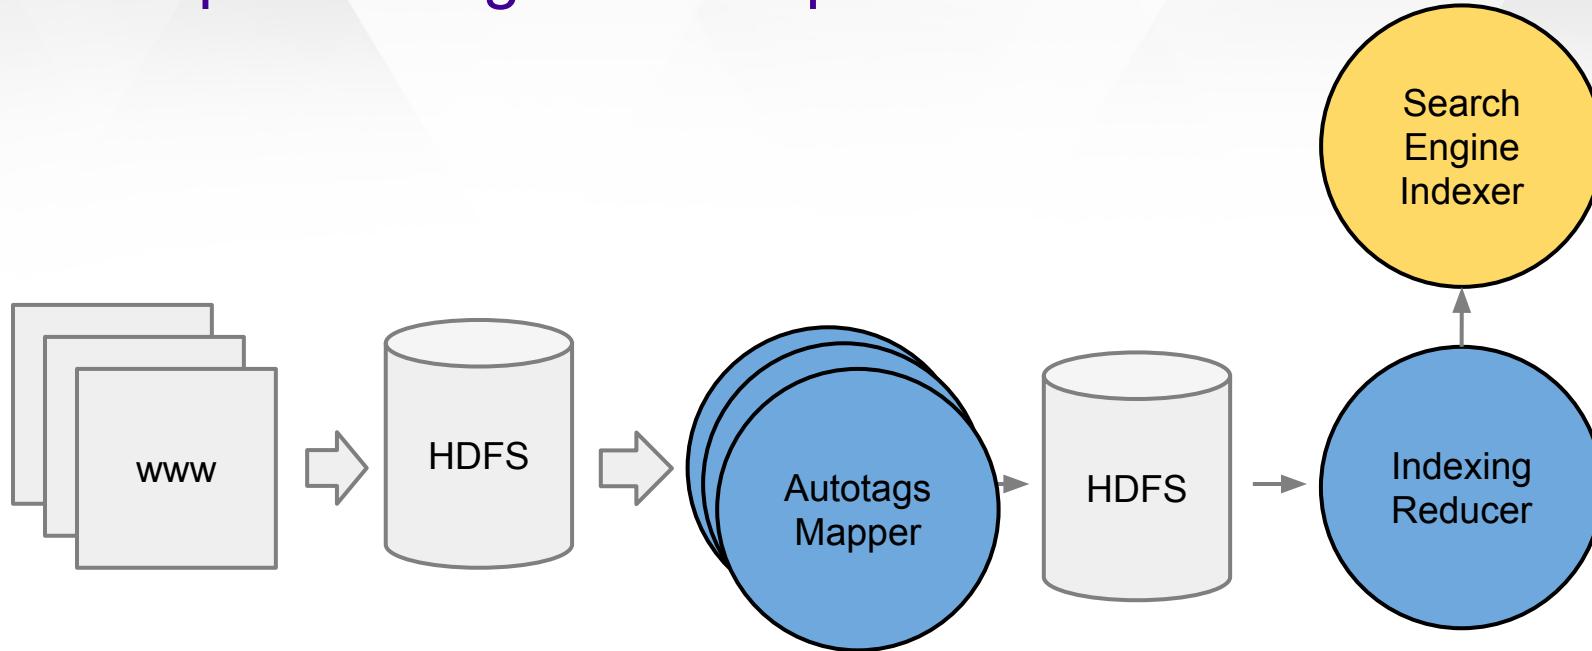
Area of Improvements

- Reducers have all training data in memory
 - A potential bottleneck if we want to expand training set for any given set of classifiers
 - Known path to avoiding this bottleneck
- Incrementally update models as new data arrives

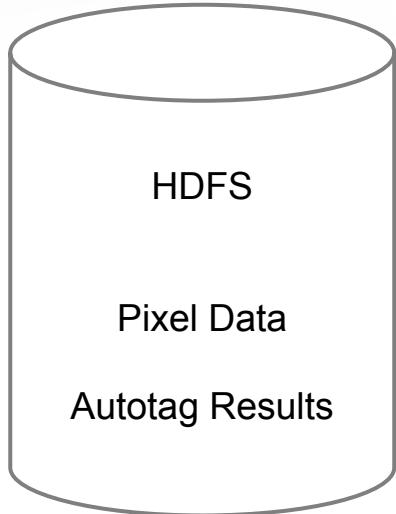
Our system for tagging the flickr corpus

- Batch processing with Hadoop
- Stream processing with Storm

Batch processing in Hadoop



Hadoop Auto tagging



Pixel Data

```
1 <base64>
2 <base64>
3 <base64>
4 <base64>
.
.
n <base64>
```

photo_id, pixel_data

Autotag Results

```
1 { "cat": .98, "dog": 0.3 }
2 { "cat": .97, "dog": 0.2 }
3 { "cat": .2, "dog": 0.3 }
4 { "cat": .5, "dog": 0.9 }
.
.
n { "cat": .90, "dog": 0.23 }
```

photo_id, results

Hadoop classification performance

- 10,000 Mappers
- Processed in about 1 week entire corpus

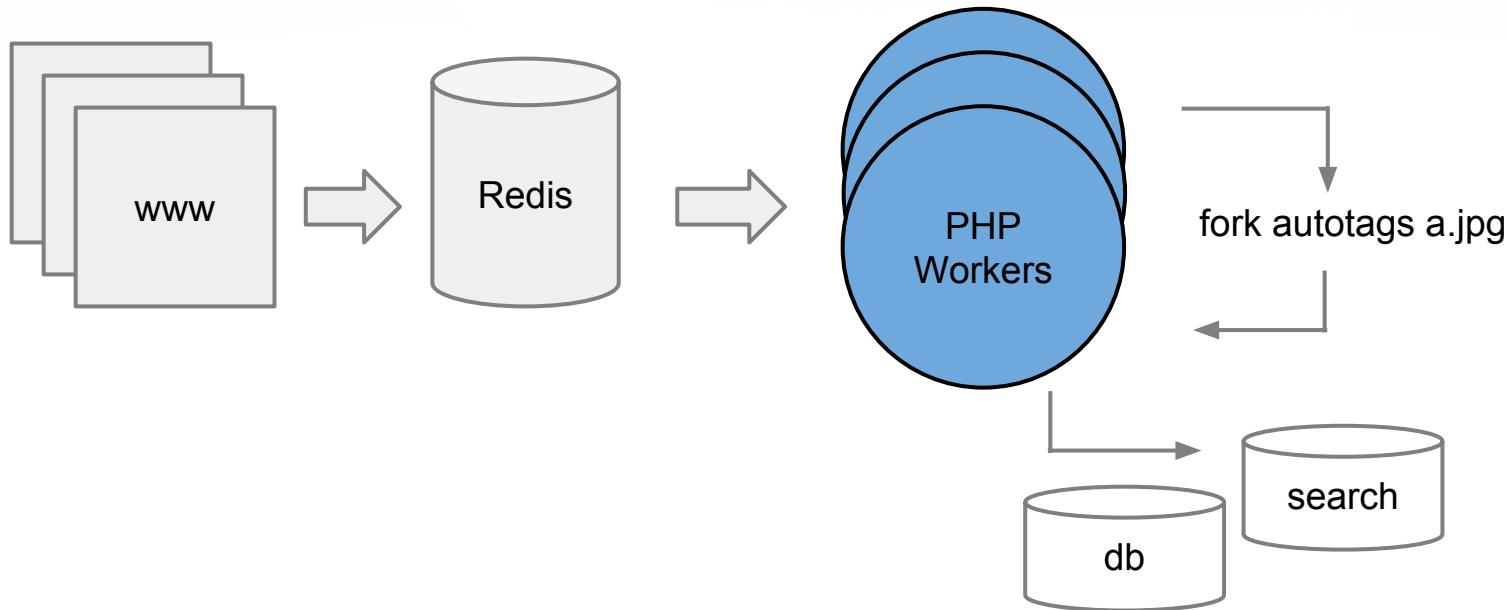
Main challenge: packaging our code

- Heterogenous Python & C++ codebase
- How do we get all of the shared library dependencies into a heterogenous system distributed across many machines?

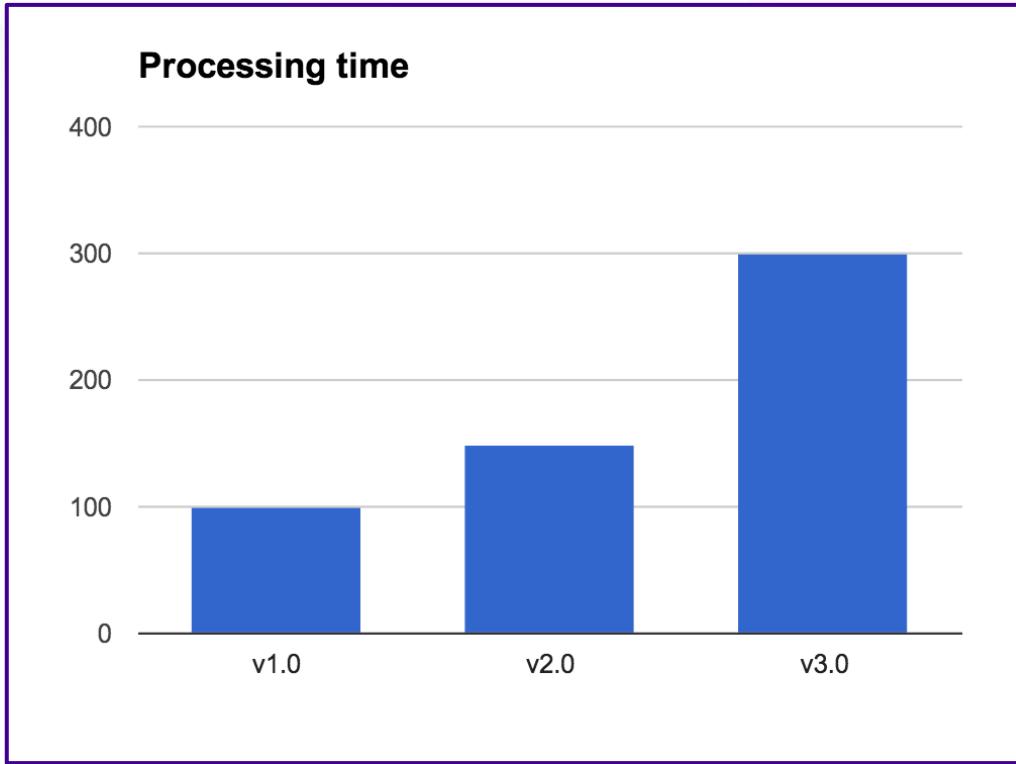
Main challenge: packaging our code

- tried: pyinstaller, didn't work
- rolled our own solution using strace
 - /autotags_package
 - /lib
 - /lib/python2.7/site-packages
 - /bin
 - /app

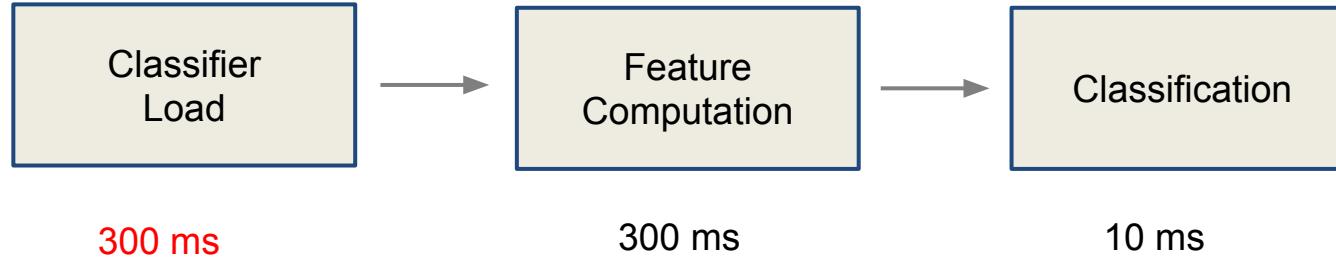
Stream processing (our first attempt)



Why we started looking into Storm



Why we started looking into Storm



Classifier loading accounted for
49% of total processing time

Solutions we evaluated at the time

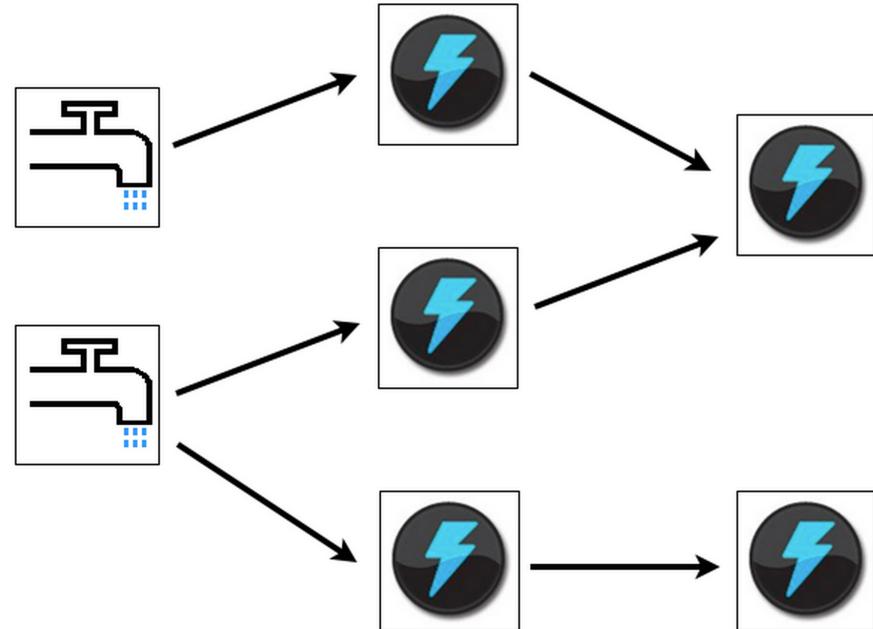
- Mem-mapping model
- Batching up operations
- **Daemonize our autotagging code**

Enter Storm

	Computational Primitives	Use case
Hadoop	Map & Reduce	Batch Processing
Storm	Spout & Bolts	Stream Processing

Storm computational framework

- Spouts
 - The source of your data stream. Spouts “emit” tuples
- Tuples
 - An atomic unit of work
- Bolts
 - Small programs that processes your tuple stream
- Topology
 - A graph of bolts and spouts



Defining a topology

```
TopologyBuilder builder = new TopologyBuilder();

builder.setSpout("images", new ImageSpout(), 10);

builder.setBolt("autotag", new AutoTagBolt(), 3)
    .shuffleGrouping("images")

builder.setBolt("dbwriter", new DBWriterBolt(), 2)
    .shuffleGrouping("autotag")
```

Defining a bolt

```
class Autotag(storm.BasicBolt):

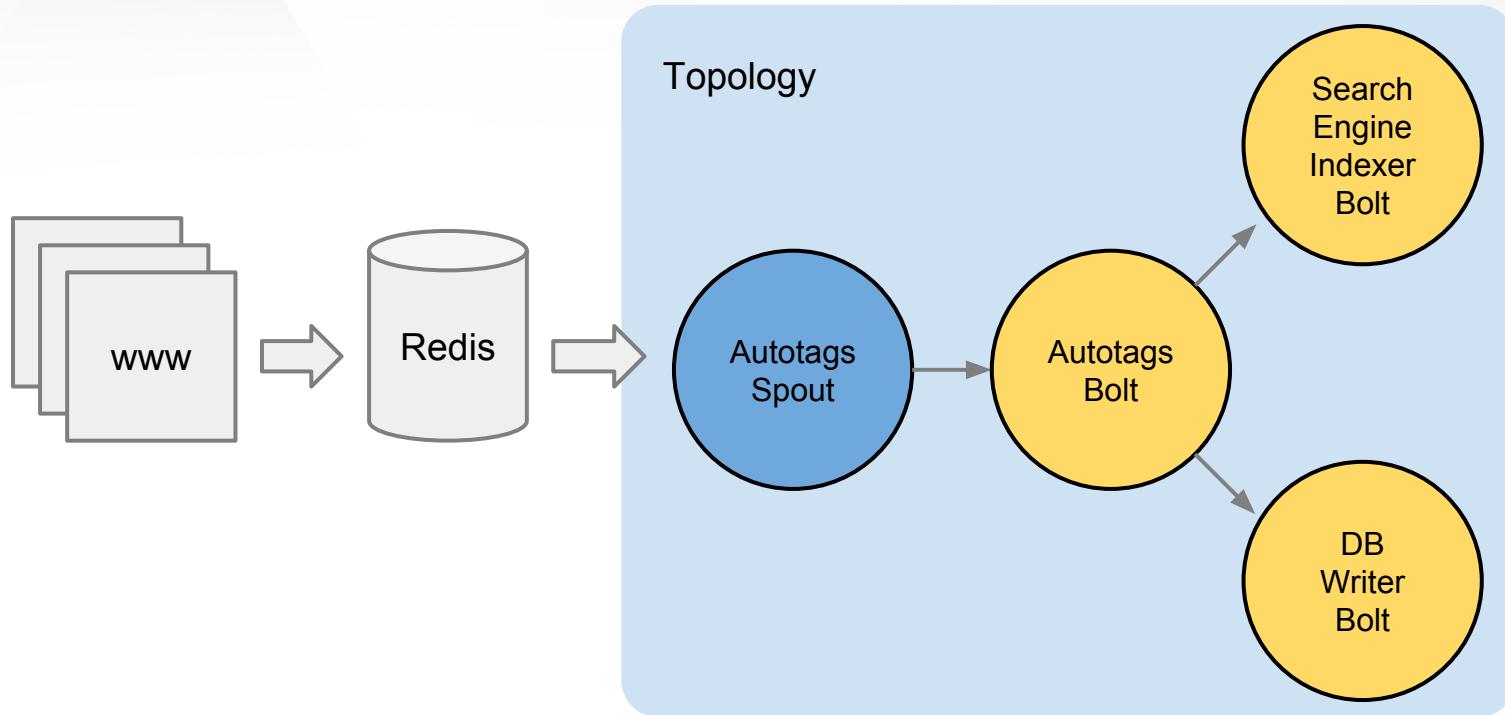
    def __init__(self):

        self.classifier.load()

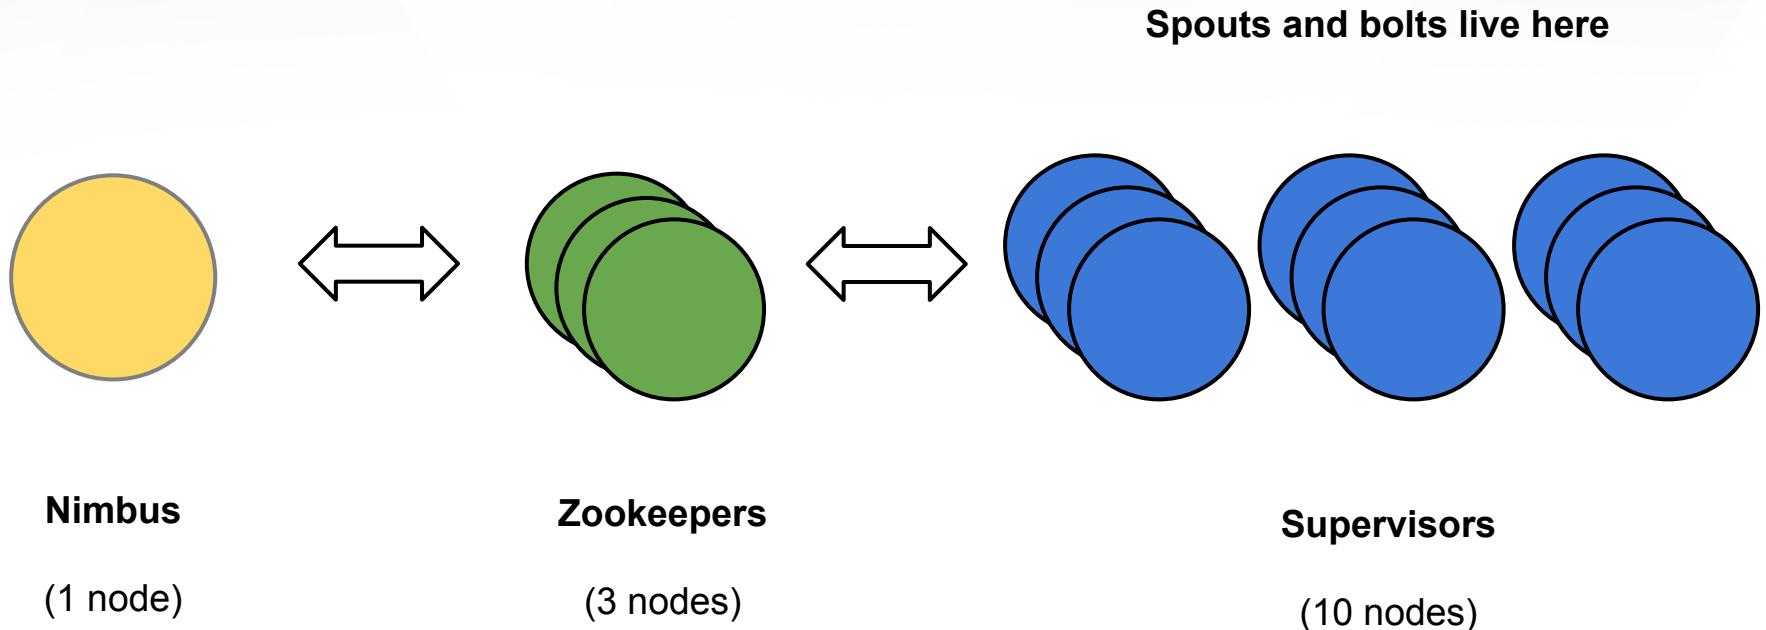
    def process(self, tuple):

        <process tuple here>
```

Our final solution



Storm architecture



Fault tolerant against node failures

Nimbus server dies

- no new jobs can be submitted
- current jobs stay running
- simply restart

Supervisor node dies

- All unprocessed jobs for that node get reassigned
- Supervisor restarted

Zookeeper node dies

- Restart node, no data loss as long as not too many die at once

storm commands

```
storm jar topology-jar-path class
```

storm commands

storm list

storm commands

```
storm kill topology-name
```

storm commands

```
storm deactivate topology-jar-path class
```

storm commands

```
storm activate topology-name
```

Is Storm just another queuing system?

Yes, and more

- A framework for separating application design from scaling concerns
- Has a great deployment story
 - atomic start, kill, deactivate for free
 - no more rsync
- No single point of failure
- Wide industry adoption
- Is programming language agnostic

Try storm out!

<https://github.com/apache/incubator-storm/tree/master/examples/storm-starter>

<https://github.com/nathanmarz/storm-deploy>

Thank you